Toward Sustainable Data Centers: Integrating Algorithmic, Architectural, and Systemic Strategies for Energy Efficiency

[1]Thejas Madhu, [2]Athul KR,[3] Neville R, [4]Sreya KA, [5]AV Aakarsh, [6]Savyo S, [7]Aryananda, [8]Vargheese Manadan

Department of Data Science
Bharata Mata College(Autonomous), Thrikkakara
thejasmadhu19@gmail.com

ABSTRACT

The rapid growth of artificial intelligence and cloud computing has increased the energy and water use of data centers, creating major environmental issues. By 2030, data centers are expected to account for 3% of global electricity consumption, with significant water needs for cooling. For instance, each query to ChatGPT uses about 519 milliliters of water. This paper offers a detailed look at various computing and infrastructure methods to improve sustainability in data centers. We discuss algorithmic innovations such as model pruning, which includes magnitude-based, structured, and random approaches which reduce electricity consumption therefore reducing the amount of heat it produces in turn reducing water usage for cooling. The paper shows on ResNet architectures that magnitude-based pruning can reduce parameters by 48.32% while also improving accuracy. The study also explores efficient model architectures used by Deepseek like Mixture of Experts (MoE) and Multi-Head Latent Attention (MLA). These enable sparse activation and latent compression, which significantly cut down inference costs. The paper demonstrates that when combined with Dual Pipeline Parallelism, these techniques can achieve up to 92% greater energy efficiency compared to traditional models.We also look at system-level methods, including server virtualization. This process consolidates workloads onto fewer physical servers to cut down on idle power use and cooling costs. It works alongside the carbon-aware MAIZX scheduling framework, which dynamically directs computing tasks to data centers or servers with the lowest carbon intensity at any given time. This approach matches compute work with the availability of renewable energy. Additionally, we explore Green Federated Learning, a new approach that optimizes decentralized model training across edge devices.By combining algorithmic improvements, architectural developments, and resource-aware system designs, this work creates a comprehensive framework for promoting sustainable computing practices in data centers. This supports global climate goals and responsible management of energy and water resources.

Keywords: Sustainable Computing, Data Centers, Model Pruning, Mixture of Experts (MoE), Multi-Head Latent Attention (MLA),Dual Pipeline Algorithm, MAIZX scheduling, Server Virtualization, Energy Efficiency, Green Federated Learning.

INTRODUCTION

The rapid growth of artificial intelligence and Machine Learning has led to a sharp rise in energy and water use in data centers, creating significant environmental issues. By 2030, data centers are expected to represent about 3% of global electricity consumption, with considerable water needs for cooling. For instance, each ChatGPT query uses an estimated 519 milliliters of water. Tackling this sustainability crisis requires a complete approach that covers algorithmic, architectural, and systemic levels of computing.

Current research often looks at these issues separately, focusing either on hardware efficiency, model improvement, or resource management. However, the biggest opportunity for change lies in combining different strategies across the whole computational pipeline. This paper presents a comprehensive framework for sustainable data centers that includes things like:
Algorithmic innovations like model pruning, Mixture of Experts (MoE), and Multi-Head Latent Attention (MLA) to cut down on computational load and memory use.
Architectural improvements such as Dual Pipeline Parallelism and server virtualization to make better use of hardware and consolidate resources.
System strategies like the carbon-aware MAIZX scheduling framework and Green Federated Learning, which optimize workload placement and decentralized training to reduce carbon and water use.

The paper provides evidence, through magnitude-based pruning that reduces ResNet parameters by 48.32% while keeping accuracy intact, that these methods together lower energy consumption, cooling needs, and related water use. A comparative analysis helps select techniques based on the application context, infrastructure limits, and sustainability goals.

This work emphasizes the importance of integrating different layers to achieve real sustainability improvements and points out key research gaps. These include the need for standard metrics, real-world validation, and flexible systems. By linking algorithmic efficiency with overall resource awareness, we present a practical pathway toward greener and more sustainable data center operations that support global climate goals.

LITERATURE REVIEW

The rising energy and water needs of modern computing, driven by the rapid growth of Artificial Intelligence (AI) and cloud services, call for a thorough approach to sustainability. Data centers are expected to make up 3% of global electricity use by 2030. Additionally, the environmental costs, such as the estimated 519 milliliters of water used per ChatGPT query, highlight the need for integrating algorithmic, architectural, and systemic strategies to manage resources responsibly[1][2]. Research in sustainable computing generally falls into three areas: algorithmic efficiencies like model pruning and sparse architectures, architectural improvements such as parallelism and virtualization, and systemic strategies involving carbon-aware scheduling and Green Federated Learning (FL)[8].
Algorithmic advancements aim to cut down on the computational load and memory use in AI models. Model Pruning includes magnitude-based, structured, and random methods that reduce electricity consumption by eliminating less important weights. For example, magnitude-based pruning has been shown to decrease ResNet parameters by 48.32% while also improving accuracy, showing that simplifying can lead to better model performance[3]. For larger models, Mixture of Experts (MoE) models offer efficiency by activating only the specific networks, or "experts," that match a given input. This reduces the computational load and energy needs compared to using one large, monolithic model[5]. A gating network manages this selective activation and helps avoid the energy waste of running unnecessary components, allowing high-performance AI with a lower carbon footprint. In addition to MoE, the Multi-Head Latent Attention MLA mechanism, used in models like DeepSeek-V3, boosts efficiency during inference by reducing Key (K) and Value (V) matrices into a smaller latent space. This compression cuts down on memory bandwidth demand and increases the Operational Intensity (OI) of the workload, which is important for saving energy on bandwidth-limited hardware[3][5].

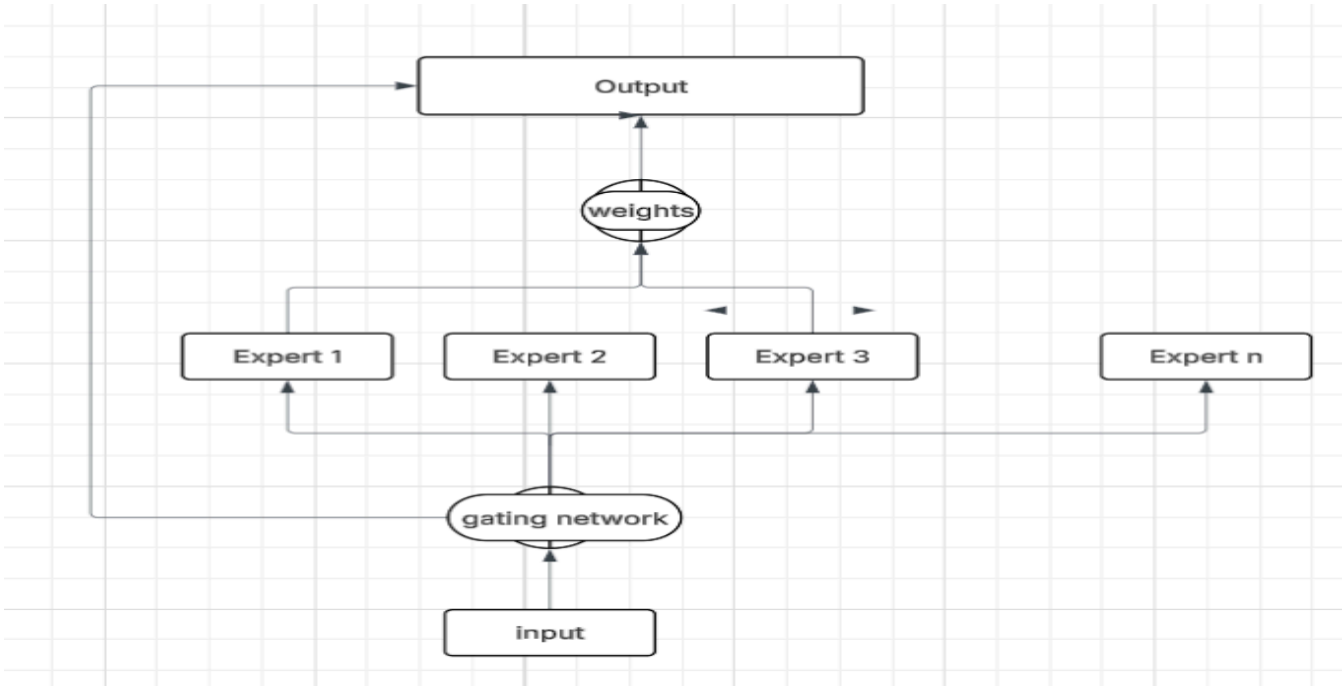


Figure 1: Mixture of experts architecture

Architectural and systemic strategies work on improving hardware use and workload placement. Server Virtualization allows for consolidating workloads, such as Virtual IEDs(vIEDs) in digital substations, onto fewer physical servers[6]. This reduces idle power use and cooling demands. For complicated distributed training tasks, the Dual Pipeline Parallelism (DualPipe) algorithm boosts computational throughput by splitting training batches into microbatches[4][5]. It uses optimizations like communication-computation overlap and bidirectional forward-backward overlap to lower idle time, or "pipeline bubbles," between processing stages.

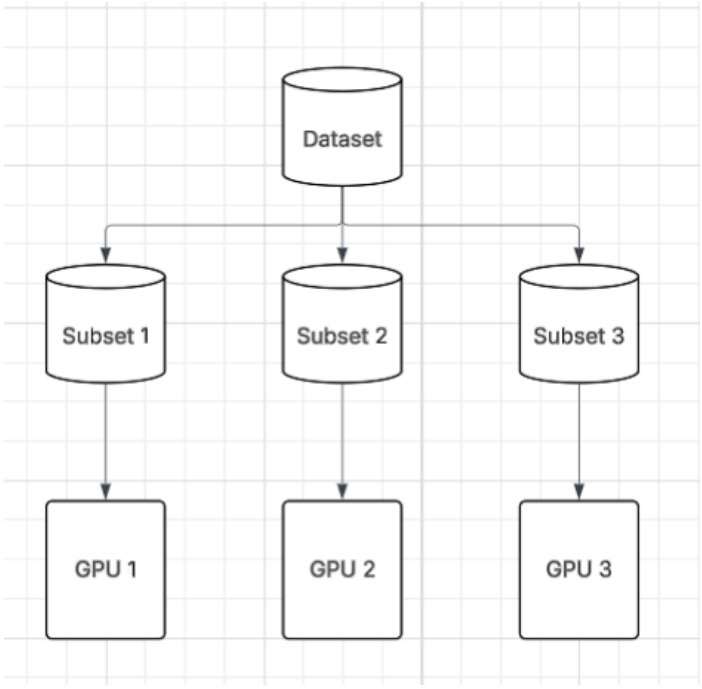


Figure 2: Data Parallelism architecture

On a larger scale, systemic strategies focus on cutting carbon emissions. The MAIZX scheduling framework uses a hybrid approach where distributed agents gather real-time power use and carbon intensity data. Its ranking algorithm assigns workloads to nodes with the lowest carbon intensity, effectively linking compute work with renewable energy availability[6][7]. Likewise, Green Federated Learning(Green FL) includes environmental sustainability in decentralized training, treating carbon footprint as an evaluation factor along with accuracy and convergence speed. Studies on Green FL indicate that the carbon emissions of an FL task are closely tied to the product of concurrency and training duration. Finding the right hyperparameters to shorten training time while minimizing communication overhead, which can make up around 50% of FL's carbon footprint, are important strategies for reducing environmental impact[8].

$$\text{MAIZ_RANKING} = w_1\,\text{CFP} + w_2\,\text{FCFP} + w_3\,\text{CP_RATIO} + w_4\,\text{SCHEDULE_WEIGHT}\,.$$

Figure 3: MAIZ Ranking Formula

While these combined approaches show great promise with methods like MoE and Dual Pipeline Parallelism achieving up to 92% greater energy efficiency compared to traditional models, challenges remain[5]. Current research lacks standardized metrics to evaluate sustainability across the full lifecycle. Most findings also need real-world validation beyond controlled settings. Achieving complete sustainability ultimately requires an integrated optimization framework that combines these algorithmic, architectural, and systemic methods based on real-time factors like workload, hardware, and environmental conditions[2].

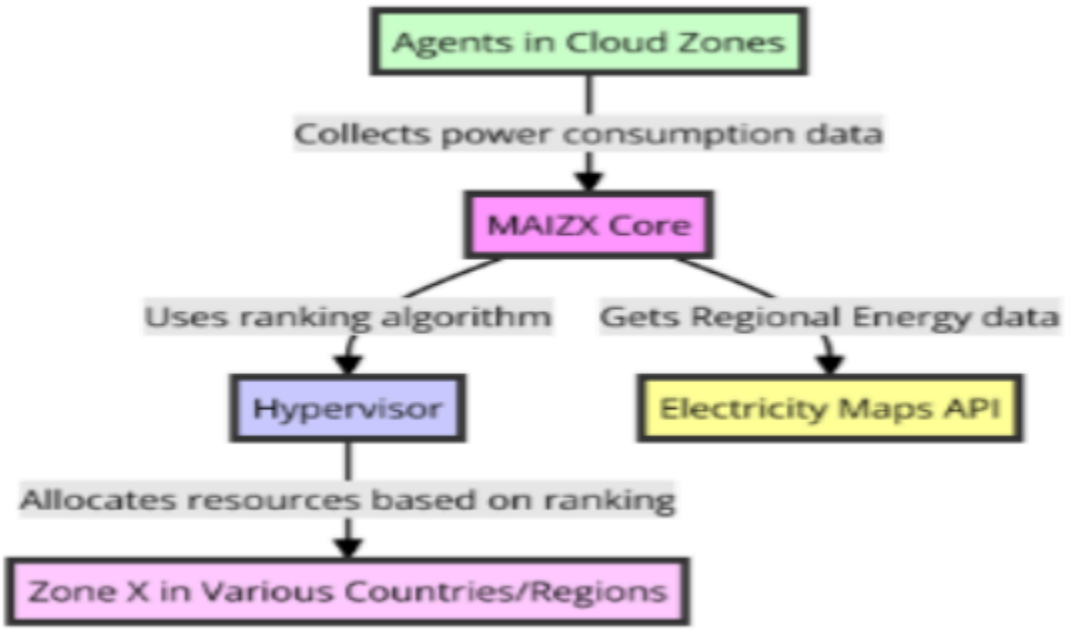


Figure 4: MAIZX framework architecture

COMPARATIVE ANALYSIS

For Artificial Intelligence or Machine Learning tasks, efficiency begins with the model architecture. Mixture of Experts (MoE) is best for large-scale, multi-task inference as it activates only the relevant model segments for each query.

To refine models, magnitude-based pruning is ideal for building accurate, compact models for edge or API deployment. Structured pruning works well in hardware-limited settings like embedded systems, trading some accuracy for quicker execution.

To train these models efficiently, Dual Pipeline Parallelism is crucial. It allows for distributed training of large models across GPU clusters, maximizing hardware use and cutting down training energy. For general data center improvement, systematic strategies are key. Server Virtualization is essential for consolidating traditional applications. This reduces the number of physical servers and lowers power consumption. The MAIZX carbon-aware framework excels in geographically distributed clouds, intelligently directing delay-tolerant workloads, such as batch processing, to locations and times with the greenest energy. This helps reduce Scope 2 emissions.

For decentralized computing, Green Federated Learning (FL) is the best choice for privacy-sensitive, cross-device training, such as in healthcare or mobile apps. It cuts down the carbon footprint of decentralized training by optimizing communication and using idle device resources.

These methods have been suggested to be effective in reducing the amount of carbon required to train, analyze and do the work these data centers are required of. Its primary way is by reducing energy consumption which is sourced from fossil fuels. The reduced consumption in energy means less heat is produced by the data centers. When less heat is produced the amount of water required to cool down these systems are also reduced meaning these methods seek to rectify both issues.

No single approach is the best. The greatest impact comes from using these methods together. Sustainable computing requires collaboration across algorithmic, architectural, and infrastructural layers.

MATERIALS AND METHODS

The research outlines a clear method for improving sustainable computing by combining algorithmic, architectural, and systemic strategies in data centers. The methods and materials focus on improving model efficiency, hardware use, and awareness of energy sources.

At the algorithmic level, three main methods aim to lower computational load and memory usage: model pruning, Mixture of Experts (MoE), and Multi-Head Latent Attention (MLA). Model pruning removes unnecessary weights to cut electricity use and heat production, which lowers cooling needs. Specifically, magnitude-based pruning targets weights with the smallest absolute values and was applied to ResNet architectures. This approach reduced parameters by 48.32% while improving accuracy. Structured pruning was applied to convolutional filters in ResNet-18 and VGG-16. It removes whole components to reduce network complexity and may lead to faster inference times, though it often needs fine-tuning to recover any lost accuracy. Random pruning served as a

comparison baseline. All pruning methods included a fine-tuning phase of an extra 10 epochs, using a reduced learning rate of 0.01 to help the sparse models recover lost performance. The Mixture of Experts(MoE) technique employs a gating network to direct queries to a relevant subset of smaller, specialized networks, known as "experts." This approach ensures only the necessary models are activated, reducing the computational load and energy requirements compared to running a single large model. Lastly, Multi-Head Latent Attention(MLA) projects Query, Key, and Value matrices into a compact, low-dimensional latent space to significantly lessen the size of the Key-Value (KV) cache. This compression lowers memory bandwidth demands and increases the Operational Intensity (OI) of the workload, improving compute utilization and energy efficiency.

At the architectural and systemic level, methods include parallel processing frameworks, virtualization, and carbon-aware scheduling. The Dual Pipeline Parallelism algorithm maximizes computational throughput and energy efficiency during the distributed training of large models. It does this by splitting training batches into microbatches and using optimizations like Communication-Computation Overlap which separates and schedules communication and compute tasks to keep devices busy and Bidirectional Forward-Backward Overlap which mixes forward and backward passes at the microbatch level to shorten overall training time. Server Virtualization consolidates workloads, such as Virtual IEDs, onto fewer physical servers, which cuts down on dedicated hardware, power use, and cooling costs. The MAIZX framework functions as a carbon-aware scheduling system with a hybrid architecture. This system uses distributed agents to gather real-time data on power consumption and carbon intensity at each node and core. The MAIZX framework, uses this data to dynamically assign workloads to nodes with the lowest carbon footprints. This balances environmental impact, performance, and operational needs through adjustable weights. Moreover, Green Federated Learning optimizes decentralized training by treating carbon footprint as a key evaluation metric. Important strategies include tuning hyperparameters like local epochs and batch size and minimizing communication overhead through model compression and quantization, which can lower data transmission volume . A predictive model is also proposed to estimate emissions based on the product of concurrency and training duration, helping practitioners choose configurations that minimize carbon footprint before they set up deployment.

TABLE 1. Pruning results for ResNet

Pruning	Sparsity	Accuracy	Training	Inference	Par Reduction (%)
Baseline	0.00	18.20	4.08	0.20	0.00
magnitude	0.10	19.60	1.00	0.20	7.11
magnitude	0.30	19.60	1.03	0.20	21.25
magnitude	0.50	20.00	1.01	0.22	35.11
magnitude	0.70	21.40	1.10	0.31	48.32
structured	0.10	19.20	1.50	0.33	7.63
structured	0.30	17.20	1.05	0.20	24.05
structured	0.50	16.80	1.01	0.21	42.26
structured	0.70	13.40	1.02	0.22	62.05
random	0.10	20.60	1.00	0.21	6.25
random	0.30	19.80	1.02	0.20	18.70
random	0.50	17.00	1.01	0.21	31.13
random	0.70	12.00	1.03	0.22	43.52

Figure 5: Pruning results using ResNet

RESULTS AND DISCUSSIONS

The results clearly show that using efficiency strategies across algorithmic, architectural, and systemic layers leads to real reductions in energy use, heat generation, and water consumption. On the algorithmic side, magnitude-based pruning was very effective, achieving a 48.32% reduction in parameters in ResNet architectures while also improving accuracy to 21.40% at 70% sparsity. This indicates that removing less important weights not only simplifies the model but can also boost performance. In contrast, structured pruning caused accuracy to decline as sparsity increased, dropping to 13.40% at 70% sparsity. This suggests that removing entire structural elements, like filters, can hurt the model's learning ability. Regardless of the pruning technique used, training times for pruned models were consistently shorter than the baseline, confirming improvements in computational efficiency.

Moreover, architectural changes for distributed training, such as combining the Mixture of Experts (MoE) model with the Dual Pipeline Parallelism (DualPipe) algorithm, showed significant efficiency. This approach achieved up to 92% greater energy efficiency compared to traditional models. The MoE model only activates specific "experts" needed for a given query, which lowers the computational load. Meanwhile, DualPipe maximizes GPU usage by reducing idle times through techniques like overlapping communication and computation and mixing forward and backward passes.

Exploring hardware-focused architectural solutions highlights the efficiency gains during model inference. The Multi-Head Latent Attention (MLA) mechanism plays a key role in lowering energy use. It does this by compressing Key (K) and Value (V) tensors into a smaller latent space, which results in a much smaller Key-Value (KV) cache. This compression significantly decreases the memory bandwidth requirement while increasing the workload’s Operational Intensity (OI), enhancing energy efficiency on hardware with bandwidth limitations.

The systemic results show how matching computing tasks with sustainable energy sources can have a practical effect. The MAIZX framework, a carbon-aware scheduling system, uses a ranking algorithm that relies on real-time data to dynamically assign workloads to nodes with the least carbon intensity. This approach prioritizes environmental impact while balancing performance and operational demands. Research on decentralized computing, particularly Green Federated Learning(Green FL), reveals that carbon emissions primarily result from the product of concurrency and training duration. Strategies to minimize this impact include optimizing hyperparameters like local epochs and batch size and reducing communication overhead. This overhead accounts for roughly 50% of FL’s carbon footprint. Techniques like int8 quantization can cut data transmission volume by up to four times.

Despite these measurable results, it is important to address ongoing research gaps and the need for integration. The conclusion stresses that no single solution is adequate, and the best outcomes arise from combining these approaches. Nonetheless, researchers face significant challenges, including the absence of clear and standardized metrics to evaluate sustainability throughout the system's lifecycle. Additionally, most favorable results, particularly those related to pruning, come from controlled experiments rather than large-scale production environments, meaning real-world confirmation of energy, water, and carbon savings is still limited. Achieving true sustainability requires overcoming these technical challenges, integrating frameworks like MAIZX with model-level optimizations, and developing adaptive systems that can dynamically choose the best techniques based on real-time workload and environmental conditions.

APPLICATIONS

The progress in sustainable computing spans different sectors, mainly focusing on large-scale AI deployment, data center infrastructure, and decentralized computing environments.

In large-scale AI training and inference, optimizing architecture and algorithms is vital for handling huge computational demands. Mixture of Experts (MoE) models are especially useful for large-scale, multi-task inference because they use a gating network. This network activates only the relevant specialized model segments, or "experts," for each query. As a result, it greatly reduces the computational load and energy needs compared to operating one large model. During the development of these models, the Dual Pipeline Parallelism algorithm is used for distributed training across GPU clusters. This strategy maximizes hardware use and cuts training energy by reducing idle time through methods like Communication-Computation Overlap and Bidirectional Forward-Backward Overlap. Additionally, the Multi-Head Latent Attention(MLA) mechanism is applied to large language models during inference. It saves processing costs by compressing Key (K) and Value (V) tensors into a smaller, low-dimensional latent space. This compression shrinks the Key-Value (KV) cache size and lowers memory bandwidth demand, which is key for energy efficiency.

In terms of data center and cloud infrastructure management, strategies focus on resource consolidation and efficient workload placement. Server Virtualization is crucial for consolidating traditional applications and workloads onto fewer physical servers. This approach lowers the total number of hardware devices, which reduces overall power consumption and cooling needs. One example is the use of Virtual IEDs in substations, where protection and control functions are centralized onto fewer servers. To lessen environmental impact, the MAIZX carbon-aware framework is designed for geographically distributed clouds. Its ranking algorithm smartly directs delay-tolerant workloads like batch processing to servers in areas or at times when electricity has the lowest carbon intensity. This dynamic scheduling helps cut down Scope 2 emissions by aligning compute work with the most available renewable energy.

In edge and decentralized computing, sustainability efforts aim to optimize distributed operations. Green Federated Learning (FL) fits well for privacy-sensitive, cross-device training, especially in fields like healthcare or mobile apps. The goal is to lower the carbon footprint tied to decentralized training. This is done by optimizing hyperparameters and reducing communication overhead, which makes up about 50% of FL's emissions. Techniques like int8 quantization can decrease data transmission volumes by up to four times. Lastly, model pruning techniques help create compact and efficient models for various uses. Magnitude-based pruning is effective for building compact models for edge or API deployment, achieving considerable parameter reduction while maintaining or improving accuracy. Structured pruning is useful in hardware-limited settings like embedded systems, where removing entire structural components speeds up execution, even if it comes with a slight accuracy trade-off.

CONCLUSION

The pursuit of sustainable data center operations relies on using strategies across algorithmic, architectural, and systemic levels. The research shows that combining techniques such as model pruning, sparse architectures like Mixture of Experts (MoE), and system-level tools like carbon-aware scheduling and virtualization can lead to significant reductions in energy consumption, cooling needs, and water use. The findings, such as magnitude-based pruning cutting ResNet parameters by 48.32% while also improving accuracy, demonstrate how algorithmic efficiency can enhance performance and reduce complexity. Additionally, architectural innovations like MoE, when paired with Dual Pipeline Parallelism, can achieve up to 92% greater energy efficiency than traditional models. This proves that maximizing hardware use through specialized parallelism results in substantial energy savings. Systemically, frameworks like MAIZX provide a way to lower environmental impact by dynamically assigning

workloads based on the real-time availability of renewable energy, matching computing tasks with the lowest carbon intensity.

Despite these encouraging results, fully realizing a sustainable computing infrastructure needs to address several key challenges and research gaps. A significant hurdle is the lack of clear and standardized metrics to evaluate sustainability across the entire computational pipeline, from model training and inference to cooling and hardware production. This gap limits fair comparison and thorough evaluation of different methods. Moreover, many positive results, such as efficiency gains from model pruning, come from controlled experiments rather than large-scale production settings, meaning actual energy, water, and carbon savings still need real-world validation. Important trade-offs, like those between computational efficiency and model accuracy, or the financial feasibility of specific architectural changes, also require further investigation.

Moving forward, future research should focus on creating integrated optimization frameworks that can combine techniques like MoE routing, MLA compression, and carbon-aware scheduling within a single automated pipeline. This flexible framework should intelligently select the best methods based on real-time factors like workload, hardware status, and environmental conditions, going beyond theoretical benefits to deliver measurable environmental results in live data centers. Establishing a standardized benchmarking suite and certification protocol is also crucial to equip the industry with the tools needed to evaluate and validate sustainability claims rigorously. Ultimately, achieving a sustainable future for computing infrastructure requires a collaborative and interdisciplinary approach to ensure that technology continues to grow in line with global climate and resource preservation goals.

References and Citations

- [1] Energy-Efficient Algorithms and Systems by Rohan Reddy, VIT-Vellore, 2024
- [2] Sustainable computing across datacenters: A review of enabling models and techniques by Muhammad Zakarya a b, Ayaz Ali Khan c, Mohammed Reza Chalak Qazani a, Hashim Ali b, Mahmood Al-Bahri a, Atta Ur Rehman Khan d, Ahmad Ali e, Rahim Khan b
- [3] OPTIMIZING MODEL PRUNING FOR ENERGY-EFFICIENT DEEP LEARNING SAIDJON KAMOLOV
- [4] Hardware-Centric Analysis of DeepSeek's Multi-Head Latent Attention by Robin Geens, Marian Verhelst
- [5] DeepSeek-V3 Technical Report by Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Daya Guo, Dejian Yang, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Haowei Zhang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Li, Hui Qu, J.L. Cai, Jian Liang, Jianzhong Guo, Jiaqi Ni, Jiashi Li, Jiawei Wang, Jin Chen, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, Junxiao Song, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Lei Xu, Leyi Xia, Liang Zhao, Litong Wang, Liyue Zhang, Meng Li, Miaojun Wang, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Mingming Li, Ning Tian, Panpan Huang, Peiyi Wang, Peng Zhang, Qiancheng Wang, Qihao Zhu, Qinyu

Chen, Qiushi Du, R.J. Chen, R.L. Jin, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, Runxin Xu, Ruoyu Zhang, Ruyi Chen, S.S. Li, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shaoqing Wu, Shengfeng Ye, Shengfeng Ye, Shirong Ma, Shiyu Wang, Shuang Zhou, Shuiping Yu, Shunfeng Zhou, Shuting Pan, T. Wang, Tao Yun, Tian Pei, Tianyu Sun, W.L. Xiao, Wangding Zeng , Wanjia Zhao, Wei An, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, X.Q. Li, Xiangyue Jin, Xianzu Wang, Xiao Bi, Xiaodong Liu, Xiaohan Wang, Xiaojin Shen, Xiaokang Chen, Xiaokang Zhang, Xiaosha Chen, Xiaotao Nie, Xiaowen Sun, Xiaoxiang Wang, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xingkai Yu, Xinnan Song, Xinxia Shan, Xinyi Zhou, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, Y.K. Li, Y.Q. Wang, Y.X. Wei, Y.X. Zhu, Yang Zhang, Yanhong Xu, Yanhong Xu, Yanping Huang, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Li, Yaohui Wang, Yi Yu, Yi Zheng, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Ying Tang, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yu Wu, Yuan Ou, Yuchen Zhu, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yukun Zha, Yunfan Xiong, Yunxian Ma, Yuting Yan, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Z.F. Wu, Z.Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhen Huang, Zhen Zhang, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhibin Gou, Zhicheng Ma, Zhigang Yan, Zhihong Shao, Zhipeng Xu, Zhiyu Wu, Zhongyu Zhang, Zhuoshu Li, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Ziyi Gao, Zizheng Pan

- [6]Design and Performance Assessment of a Virtualized IED for Digital Substations by ALAILTON J. ALVES JUNIOR*, DENIS V. COURY*, MEMBER, IEEE, and RICARDO A. S. FERNANDES*, SENIOR MEMBER, IEEE
- [7]MAIZX: A Carbon-Aware Framework for Optimizing Cloud Computing Emissions by Federico Ruilova, Ernst Gunnar Gran, Sven-Arne Reinemo.
- [8]Green Federated Learning by Ashkan Yousefpour 1 Shen Guo 1 Ashish Shenoy 1 Sayan Ghosh 1 Pierre Stock 1 Kiwan Maeng 1 2Schalk-Willem Kr ¨uger 1 Michael Rabbat 1 Carole-Jean Wu 1 Ilya Mironov